

Transcript for the Presentation of the Assignment: Neural Network Models for Object Recognition

INTRO PAGE / TITLES AND TOPIC

This assignment will be used for the partial fulfillment of requirements of Module 3 Machine Learning that started October 2024 as part of the PGD Certificate in Artificial Intelligence at University of Essex Online.

Executive Summary:

- The assignment aimed to train a Neural Network using the dataset of the Canadian Institute For Advanced Research-10, in short CIFAR 10. The task was to evaluate the performance of the dataset and use Python for data mining. Key aspects included data preparation, model architecture and training, design strategies as well as challenges and learnings. A Convolutional Neural Network, in short CNN, was chosen as the preferred deep learning model since it is very suitable for structured data like images, as used in this dataset.
- The problem statement revolved around exploring whether a Convolutional Neural Network can effectively be trained for object recognition using this dataset, and if so, how well the performance and accuracy in identifying objects across ten distinct categories will be.
- Key outcomes included a 55 % accuracy and a 1.24 loss, showing a good first result that is underperforming in comparison with benchmarks. The model struggled with underrepresented classes like "cat" and "bird" and quite similar pairs such as airplanes and ships.
- Working on model refinement throughout the course of an entire week improved stability, and personal programming skills, but there is still room for improvement.

Agenda

This presentation is structured into 4 parts, starting with the Introduction, continuing with the workflow of the model development and afterwards going into model optimization and validation. Learnings and other insights finalize this presentation.

Introduction: Training CNNs for Object Recognition with CIFAR-10

Artificial Neural Networks (ANNs) are a subset of AI systems, often described as functioning in a manner similar to the human brain. They are widely applied in various domains, with speech and image recognition being among the most prominent use cases, as described by Chi Leung Hui (Chi Leung Hui, P., 2023).

- According to Al-Malah, Neural networks are popular due to their capacity to learn from data and make predictions or informed decisions. Subsets of image recognition include applications that identify faces and objects and categorize images.
- Neural Networks are being applied in these areas to assist in interpreting visual data in security systems, social media, and autonomous vehicles (Al-Malah, K.I.M. 2024).
- Al-Malah continues to say that Convolutional Neural Networks (CNNs) are a type of deep learning model optimized for handling structured data like images, transforming inputs through weighted layers into 3-D outputs known as activations. They are versatile enough to be used in areas such as text, audio and signal processing, with their structure customized for specific tasks by varying the Layer configurations and quantities (Al-Malah, K.I.M. 2024).
- According to Aceves-Fernandez, a loss function is crucial for evaluating the performance of a neural network and CNN training is focused on minimizing exactly this loss (Aceves-Fernandez, M. A. 2020).
- The objective of this assignment, finally, is training a CNN for object recognition using the CIFAR-10 dataset and on this slide we see a sample of labeled images from this very same dataset. The assignment covered dataset preparation, model design, training, evaluation, and insights.

Understanding the CIFAR-10 Dataset and the split methodology

- Let's address why CIFAR-10 was chosen for this project. As Mattmann highlights, it is a widely used benchmark dataset for image classification tasks. It contains 60,000 images, each measuring 32x32 pixels, divided into 10 categories such as airplane, automobile, bird, cat, and others (Keras (n.d.) 'CIFAR-10 Dataset API'). Its manageable size makes it an ideal choice for this assignment, because it has the right balance between being complex and yet feasible for the use case of developing and testing visual recognition models.
- The common datasplit of 80/10/10 was chosen, meaning 80% for training, 10% for validation, and 10% for the test set, resulting in 48,000 training and 6,000 validation and test images each.

Witten states that a Validation Set is Used during training to adjust hyperparameters and optimize model complexity, ensuring the model does not overfit or underfit while adjusting its architecture or parameters.

He continues to say that the Test Set is Used after training to evaluate the final model's performance on unseen data, ensuring unbiased metrics and confirming generalization.

Data Exploration and Neural Network Architecture

- We continue with the second part of the presentation, which focuses on the Model Development Workflow.
- To begin, the CIFAR-10 dataset was prepared by loading the data and normalizing pixel values to a range between 0 and 1, making it easier for the model to process. The dataset was divided into training, validation, and test sets in an 80/10/10 ratio, ensuring sufficient data for each stage of training and evaluation.
- The labels were one-hot encoded to align with the model's requirements. To better understand the dataset, a sample image with its label was displayed, along with a grid view of training images, allowing to analyze the diversity of classes. Data analysis also included verifying the shapes of the data and examining class distributions to confirm that the dataset was ready for modeling.
- The neural network was designed with three convolutional layers, each using Rectified Linear Unit (ReLU) activation, batch normalization, max-pooling, and dropout. These layers worked together to extract features effectively while minimizing overfitting. A dense layer with 256 units was added next, also using ReLU activation and dropout, to further refine learning and ensure generalization. The reason the Rectified Linear Unit was used throughout the convolutional and fully connected layers because it helps the model train efficiently by preventing issues like weakening gradient signals (Chi Leung Hui, 2023; Witten et.al. 2017).
- The final layer employed softmax activation to convert the network's outputs into probabilities, predicting the likelihood of each input belonging to one of the 10 categories.
- The model was trained using the Adam optimizer, which is known for automatically adjusting learning rates, and the categorical cross-entropy loss function, which is well-suited for multi-class classification tasks (Al-Malah, 2024). During training, accuracy tracking was used to monitor the model's performance. A utility tool was also implemented to visualize images and their labels, offering deeper insights into the model's structure, important components, and overall complexity.
- The bar chart shown here illustrates how computational demands are distributed across the layers. Most of the workload is handled by the Dense and 2D Convolutional layers, which are responsible for learning features and making predictions. In contrast, pooling and dropout layers

focus on improving efficiency and reducing overfitting. Together, these components enable the Convolutional Neural Network (CNN) to effectively learn and generalize, making it highly capable of tasks like image recognition and classification.

Improving Robustness and Generalization with Augmentation

This part is about training the model.

- To make the model more robust and better at generalizing new data, real-time data augmentation techniques were applied to the training set. These included rotations, shifts, flips, zooms, and shearing, which expanded the scope of the dataset.
- Training was carried out on the augmented dataset in batches of 64 and lasted up to 35 epochs. Validation data was used throughout to monitor the model's performance and ensure it was advancing as expected.
- Early stopping was implemented to stop training when the validation loss no longer improved, which helped to prevent overfitting. Additionally, model checkpointing was implemented to save and restore the weights from the best-performing epoch. Together, these methods ensured that the model learned effectively from the data while maintaining accuracy and reliability when making predictions on new images.
- The training process required several adjustments and repetitions over the course of a week due to challenges, which will be discussed later in the presentation.
- Finally, the Training-Validation Loss Trend highlights that after multiple refinements, the model was fairly well-trained and demonstrated strong performance on unseen data.

From Metrics to Visuals: Understanding Model Behavior

We have arrived at the third section, which concerns Model Optimization and Validation. In this section we delve into performance metrics, what went well in optimizing the model, and what could have led to a better result.

Let's look at some first performance metrics of this trained model.

The model achieved 55% accuracy, meaning it correctly predicted the class for slightly more than half of the test samples. While this is better than random guessing (which is equivalent to 10% accuracy for a dataset with 10 classes), it certainly has the potential to arrive at better results.

The classification metrics results, as seen on the right-hand side of this slide, indicate that the model performs unevenly across different classes. It performs well on structured and easily distinguishable classes like **ship** (F1-score: 0.72) and **truck** (F1-score: 0.65), likely due to their clear patterns. In contrast, it struggles with less distinct or more complex categories like **cat** (F1-score: 0.32) and **bird** (F1-score: 0.35), most likely due to their more subtle patterns that make differentiation challenging.

The macro average and weighted average results confirm that while the model is effective at recognizing certain well-defined objects, it has difficulty generalizing more ambiguous categories. Possible ideas for improvement include further fine-tuning hyperparameters, or additional training with more diverse and representative samples. All this could help achieve a better performance across all classes.

Performance Overview: Consistent Accuracy with Notable Misclassifications in Similar Classes

We continue with further metrics and their evaluation.

The model achieves a test accuracy of 55.22% and a loss of 1.24, aligning with validation trends and demonstrating consistent performance across datasets. While this indicates stable generalization, it is significantly worse than the state-of-the-art accuracy that falls anywhere between 90 and 99%. The confusion matrix reveals strengths in predicting classes like frog (6) and ship (8) but highlights struggles with visually or contextually similar categories, such as airplane vs. ship (0 vs. 8) and cat vs. frog (3 vs. 6), likely due to shared visual features.

This would be another area where the model could be tweaked to achieve better results. Some of these tweaks include a more diverse and larger datasets like CIFAR-100 to improve the model's ability to differentiate between classes. Hyperparameters such as batch size, learning rate, and number of epochs could be further explored and adjusted. Also, more advanced architectures could be used or just a higher number of filters in convolutional layers for more refined outcomes.

Refining The Model: Tackling Overfitting with Data Split Adjustments

In the final section called Learning, we will be covering what was done to optimize the model's stability, accuracy, and generalization by refining certain parameters. Also, which challenges were faced and which insights were gained.

The initial model faced challenges with overfitting and low validation accuracy due to an imbalanced 70/20/10 data split, too many epochs (50), and a batch size of 64. Early fixes included adding dropout and BatchNormalization layers, reducing the batch size to 32, and limiting epochs to 30. However, these adjustments caused unstable validation loss and worsened overfitting.

To fix these issues, the data split was corrected to 80/10/10, creating a better balance between training, validation, and test sets. This adjustment laid the foundation for better performance. Next, the learning rate was reduced significantly (to 0.0003–0.0004), allowing the model to make smaller, more stable updates during training. The early stopping patience was increased to 5–7 epochs, giving the model extra time to improve without risking overfitting. Also, the Adam optimizer was selected because it updates weights efficiently, speeding up learning. Finally, random seeds were set to ensure consistent and reproducible results.

At this point it's important to mention, that certain model design choices were intentionally simplified to prioritize achieving a stable, functional model over optimizing for peak performance. This approach was primarily influenced by limited programming expertise, as implementing and troubleshooting more complex adjustments proved time-intensive. Consequently, the convolutional layers were capped at three, the architectural complexity was kept minimal, and the addition of extra features was carefully chosen.

Additional adjustments included setting the batch size back to 64, limiting training to 35 epochs, and applying data augmentation techniques (e.g. flips and rotations). These techniques helped the model perform better on new, previously unseen data.

To better understand the dataset, a tool called utility function was used to display grids of images with their labels. The model's architecture was also simplified by reducing the number of filters in the 2D convolutional layers and the number of neurons in the dense layers. These changes helped the model generalize more effectively and reduced the risk of overfitting as well.

Learnings and Conclusion

Finally, a few more insights into personal challenges but also take-aways from this assignment.

Throughout the third module of this course, I saw significant growth in my Python programming skills, although I'm still at a beginner level compared to those who use it more frequently.

When I started working on this assignment, I made an error with the test split early on, setting it at 70/20/10 instead of the desired 80/10/10. This mistake slowed progress and worsened results despite my adjustments. Once I identified and corrected the split, combined with reducing the epoch count, adding early stopping, and using the original batch size, the model became much more usable. Debugging was a key part of this process, involving careful checks and code validation throughout training, and using random seeds to ensure reproducibility.

Working directly in Python proved faster than on Cloud platforms, but it came with challenges, like slower debugging and fewer resources compared to platforms like Google Collab. However, through iterative refinements of the code and model, I was able to improve my technical skills and at the same time deepen my understanding of neural networks. This slower process gave me an understanding of how Convolutional Neural Networks (CNNs) function in real-world applications, including image recognition and prediction, and how training and optimization work.

Visualizing progress through graphs and metrics also proved helpful for my learning and helped me achieve stable, more reliable results.

A potential continuation of this project could involve a few previously mentioned features. The model's performance could be more powerful by using larger datasets like CIFAR-100, optimizing hyperparameters further such as batch size, learning rate, and epochs, and implementing advanced architectures or increasing convolutional filters for improved class differentiation and refined outcomes.

References

Books and Edited Volumes:

- Aceves-Fernandez, M.A. (ed.) (2020) *Advances and Applications in Deep Learning*. London: IntechOpen.
- Al-Malah, K.I.M. (2024) *Machine and Deep Learning Using MATLAB: Algorithms and Tools for Scientists and Engineers*. 1st edn. Hoboken, NJ: John Wiley & Sons.
- Chi Leung Hui, P. (2023) *Artificial Neural Networks: Recent Advances, New Perspectives and Applications*. Available at: <https://www.intechopen.com> (Accessed: 16 January 2025).
- Mattmann, C. (2021) *Machine Learning with TensorFlow*. 2nd edn. New York: Manning Publications Co. LLC.
- Witten, I.H., Frank, E., Hall, M.A., and Pal, C.J. (2017) *Data Mining: Practical Machine Learning Tools and Techniques*. 4th edn. Cambridge, MA: Morgan Kaufmann.

Articles:

- Jordan, K. (2024) '94% on CIFAR-10 in 3.29 Seconds on a Single GPU', *arXiv preprint*. Available at: <https://arxiv.org/abs/2404.00498> (Accessed: 16 January 2025).

Websites:

- Papers with Code (n.d.) 'Comparison Benchmarks: Image Classification on CIFAR-10'. Available at: <https://paperswithcode.com/sota/image-classification-on-cifar-10?> (Accessed: 12 January 2025).
- Scikit-learn (n.d.) 'Supervised Neural Networks'. Available at: https://scikit-learn.org/1.5/modules/neural_networks_supervised.html (Accessed: 10 January 2025).
- TensorFlow (n.d.) 'Beginner Quickstart'. Available at: <https://www.tensorflow.org/tutorials/quickstart/beginner> (Accessed: 08 January 2025).
- Kaggle (n.d.) 'CIFAR-10 Basic CNN with PyTorch'. Available at: <https://www.kaggle.com/code/icofield/cifar-10-basic-cnn-with-pytorch> (Accessed: 28 December 2024).

- Kaggle (n.d.) 'CIFAR-Conv'. Available at: <https://www.kaggle.com/code/cvvcvccvvcvc/cifar-conv> (Accessed: 21 December 2024).
- Keras (n.d.) 'CIFAR-10 Dataset API'. Available at: <https://keras.io/api/datasets/cifar10/> (Accessed: 21 December 2024).